

# Sparsifying Large Language Models

## Seminar 1

ECE 718 - Compiler Design for HPC

November 4, 2025

Tahar HERRI

# Outline

- 1 What is the sparsification of LLMs?
- 2 Why is sparsification important?
- 3 Overview of sparsification methods
- 4 Demo – Seminar 1

<https://github.com/TaharHERRI/Edu-Sparsify-LLMs>

# What is the sparsification of LLMs?

- **Sparsification** = reducing the number of active elements in a neural network (not only weights), to decrease memory and computation costs.
- This involves **removing or zeroing-out** less important components while striving to preserve model accuracy.
- What can be sparsified?
  - *Weights*: individual connections.
  - *Neurons or attention heads*: functional units.
  - *Layers or blocks*: higher-level architecture elements.
- Two complementary goals:
  - *Structural pruning*: determining which elements to remove.
  - *Sparse execution*: leveraging hardware/software to skip pruned parts and accelerate inference.

Inspired by: Han et al., "Learning both Weights and Connections," NeurIPS 2015

# Dense vs Sparse Representations

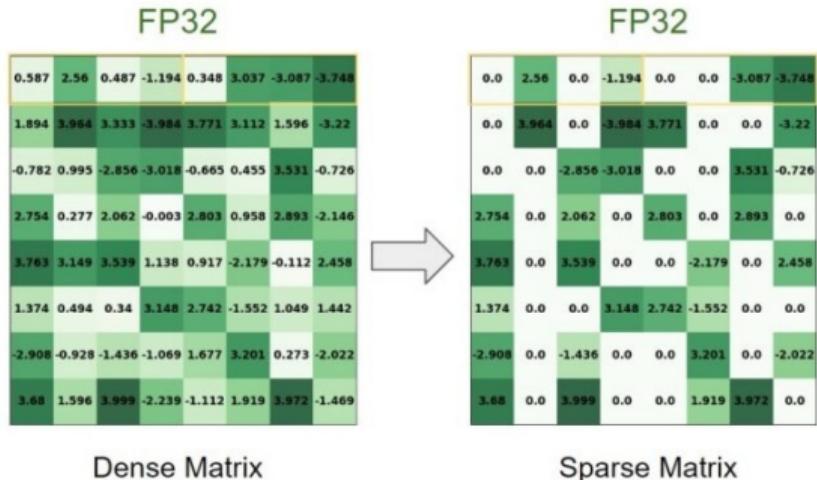


Figure: 1

Training sparsity workflow — NVIDIA (2023)

Source: NVIDIA Developer Blog

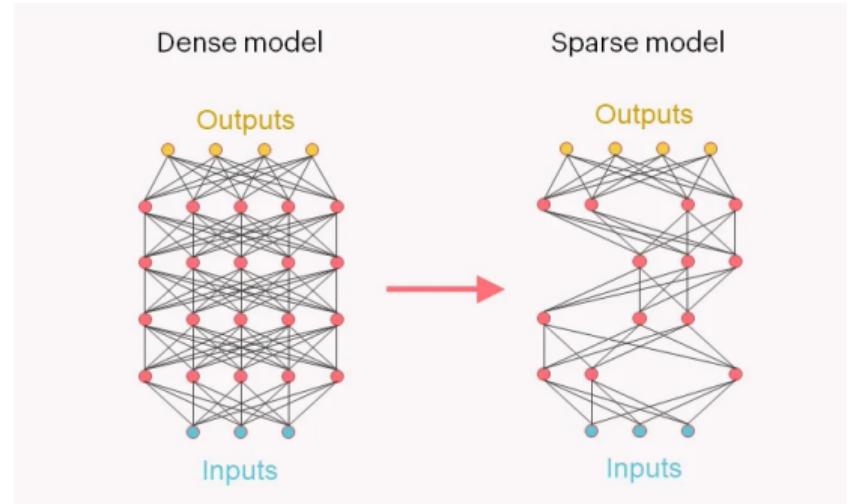


Figure: 2

Sparse LLM example — Graphcore (2021)

Source: Graphcore & Aleph Alpha

# Why is sparsification important?

- LLMs have **hundreds of billions of parameters** (GPT-4, Gemini...).
- Inference cost scales roughly linearly with parameter count.
- Sparsification enables:
  - **Lower memory usage**
  - **Faster inference** (fewer FLOPs)
  - **Lower energy cost**
  - **Deployment on edge/limited hardware**
- Complements **quantization, distillation, and LoRA (Low-Rank Adaptation)**.

Refs: Kurtic et al., 2023; Mishra et al., 2021

# Overview of Sparsification Methods

How we make large models sparse in practice

# Unstructured Pruning

- Removes individual weights (typically smallest-magnitude).
- Very flexible: high accuracy retention.
- Irregular sparsity → hard to accelerate without special kernels (CSR, ...).
- **Recent methods:**
  - **SparseGPT (2023):** fast pruning without fine-tuning
  - **Wanda (2023), SPQR (2023):** post-training, quantization-friendly

Refs: Frantar et al., 2023; Su et al., 2023; Dettmers, 2023

# Structured Pruning

- Removes **groups** of weights: neurons, channels, heads, or even layers.
- Produces a smaller **dense** model (no special execution format needed).
- Easy to run on CPUs/GPUs — widely supported.
- More aggressive pruning risks degrading accuracy.

Refs: He et al., "Channel Pruning for Accelerating Very Deep Networks," ICCV 2017.

## Semi-Structured Pruning (N:M, Block)

- Enforces patterns: e.g., **2:4 sparsity** (2 non-zero out of 4), or **block sparsity** (e.g.  $8 \times 8$ ).
- Supported by **NVIDIA Ampere/Hopper (cuSPARSELt)** and **PyTorch 2.x**.
- Trade-off: more regular = easier to accelerate, less flexible than unstructured.

*Refs:* Mishra et al., "Accelerating Sparse Deep Neural Networks," NVIDIA cuSPARSELt whitepaper, 2023.

# Dynamic Sparsity

- Sparsity changes during training over time or per input.
- Examples:
  - **RigL, SET**: sparse training with evolving masks
  - **MoE (Mixture-of-Experts)**: only a few experts used per token
  - **Top- $k$  activations**: only keep strongest activations
- Efficient but adds overhead for routing/scheduling.

Refs: Evci et al., "RigL: Efficient Sparse Training," NeurIPS 2020; Lepikhin et al., "GShard," 2021.

# Recap of Sparsification Methods

| Type            | What is removed           | How it speeds up           | HW support |
|-----------------|---------------------------|----------------------------|------------|
| Unstructured    | Individual weights        | Sparse kernels (CSR, ...)  | Limited    |
| Structured      | Neurons, heads, layers    | Smaller dense model        | Full       |
| Semi-structured | N:M or block patterns     | Specialized tensor kernels | High       |
| Dynamic         | Input/time-specific parts | Activates fewer units      | Partial    |

# Demo – Seminar 1

<https://github.com/TaharHERRI/Edu-Sparsify-LLMs>

# Steps and metrics

- ① Load baseline model
  - facebook/opt-125m on CPU
  - EleutherAI/pythia-410m on GPU
- ② Measure baseline: PPL, latency/token, model size.
- ③ Apply masked Pruning (**30%**) and Freeze then Measure.
- ④ Apply masked Pruning (**30%**), Freeze and Convert to CSR Format.
- ⑤ Apply masked Pruning (**50%**), Freeze and Convert to CSR Format.

## Metrics

Perplexity, latency per token, model size, sparsity (%).

## Expected results (30–70% sparsity)

| Setup      | Model size | Latency | Perplexity       |
|------------|------------|---------|------------------|
| Dense      | 100%       | 100%    | ref              |
| Masked 30% | ≈100%      | 100%    | ≈                |
| CSR 30%    | 85–95%     | 95–105% | ≈                |
|            |            |         |                  |
| Masked 50% | ≈100%      | 100%    | ≈                |
| CSR 50%    | 60–70%     | 90–100% | ≈                |
|            |            |         |                  |
| Masked 70% | ≈100%      | 100%    | ≈+ $\varepsilon$ |
| CSR 70%    | 40–50%     | 70–85%  | ≈+ $\varepsilon$ |

At low sparsity (30%), speedups are negligible. Real gains appear from  $\geq 50\text{--}70\%$  sparsity, where compute and memory reductions start to outweigh CSR overhead. Perplexity degradation ( $\varepsilon$ ) usually stays minor if pruning is magnitude-based.

# Demo & Analysis

Thank you and GitHub link

Thank you!

Full code, notebooks, and results:

<https://github.com/TaharHERRI/Edu-Sparsify-LLMs>